
Natural Language Processing

Introduction to NLP


Sofia Serrano
sofias6@cs.washington.edu

Credit to Yulia Tsvetkov and Noah Smith for slides


Announcements

- Academic Integrity Form is out on Canvas
- A1 is out on GitLab
 - Don't see it? Reply to this thread on Ed with your NetID:
<https://edstem.org/us/courses/32306/discussion/2365366>
- Access to lecture recordings
 - No @cs.washington.edu google account? Click through to (request) access any lecture recording sooner rather than later so that we can give you access
- Make sure you can access the course machines
 - *(if connecting from off campus)* Run [Husky OnNet VPN](#) OR first ssh into an attu machine
 - `ssh yourNetID@nlpg00.cs.washington.edu` *(nlpg00-nlpg03)*
 - Not working?
 - Not a CSE major/no CSE account? Email ugrad-adviser@cs.washington.edu to request a CSE account (include your student ID number in the email) and CC Sofia
 - Still not working? Reply to this thread on Ed so that we can help troubleshoot:
<https://edstem.org/us/courses/32306/discussion/2368995>

Following up on a question from last lecture




Credit to [Phoebe Mulcaire](#) for figures



Symbolic and Probabilistic NLP

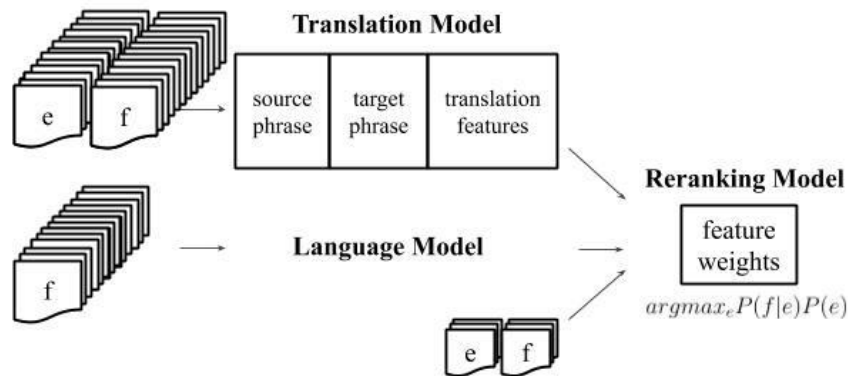
Logic-based/Rule-based NLP



~ 90s




Statistical NLP



Probabilistic and Connectionist NLP


Engineered Features/Representations



~mid 2010s



Learned Features/Representations




Linguistic Background

What does it mean to “know” a language?



(Thanks Canadian Internet Registration Authority!)



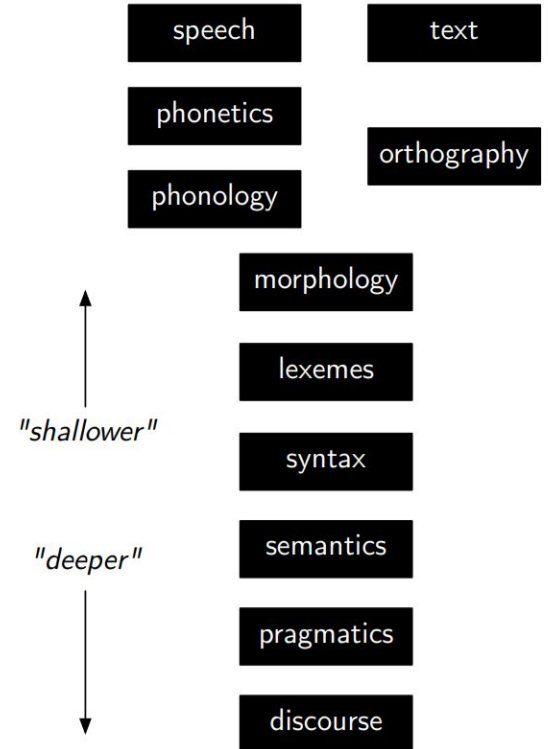
```
Last login: Mon Jan  9 08:08:57 2023 from 97.1  
[sofias6@attu8 ~]$ wc myfile.txt
```

What do we need to “tell” a computer program so that it knows more English than WC or a dictionary, maybe even as much as a three-year-old, for example?

What does an NLP system need to 'know'?

- Language consists of many levels of structure
- Humans fluently integrate all of these in producing/understanding language
- Ideally, so would a computer!

Levels of linguistic knowledge

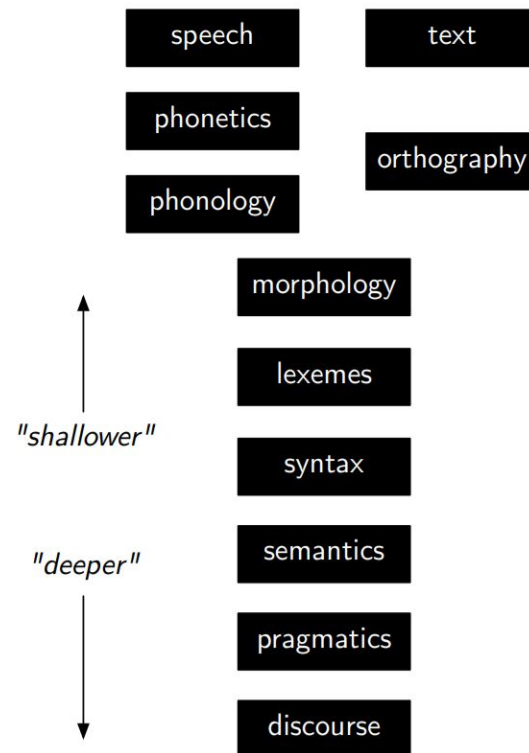


Speech, phonetics, phonology



This is a simple sentence .

/ ðɪs ɪz ə 'sɪmpl 'sɛntəns /.



Orthography


هذه جملة بسيطة

đây là một câu đơn giản

यह एक साधारण वाक्य है

This is a simple sentence .

/ ðɪs ɪz ə 'sɪmpl 'sentəns /.

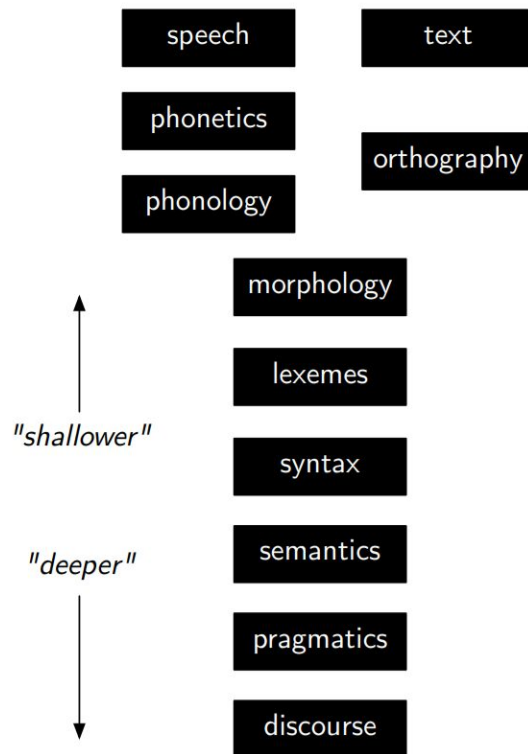


Words, morphology

- Morphological analysis
- Tokenization
- Lemmatization

Tokens This is a simple sentence .

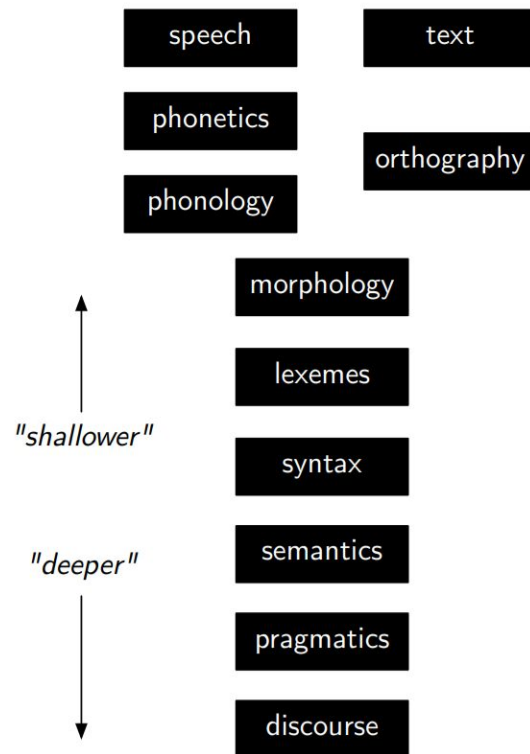
Morphology be
3sg
present



Syntax

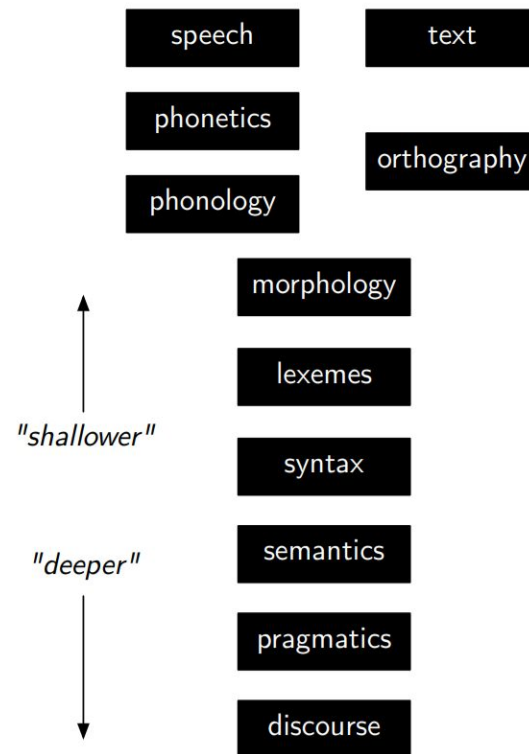
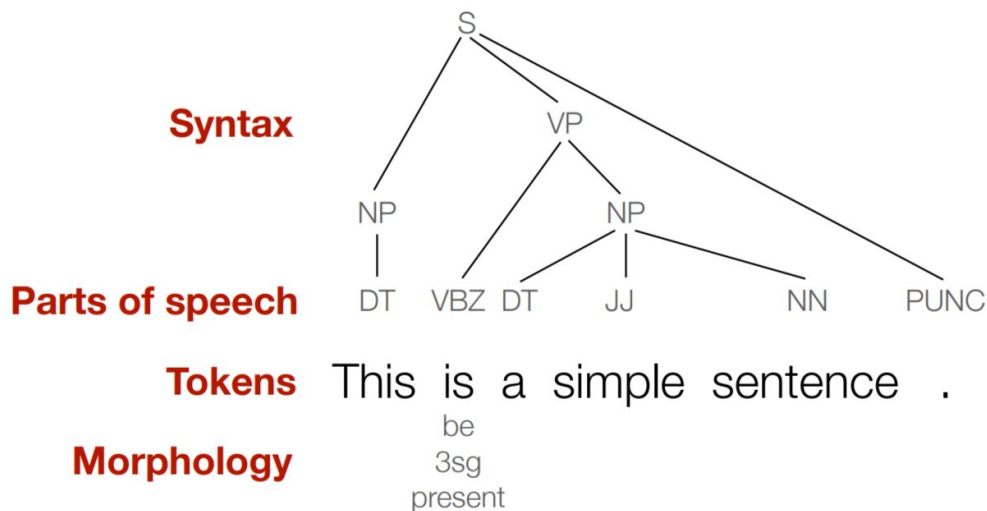
- Part-of-speech tagging

Parts of speech	DT	VBZ	DT	JJ	NN	PUNC
Tokens	This	is	a	simple	sentence	.
Morphology		be 3sg present				




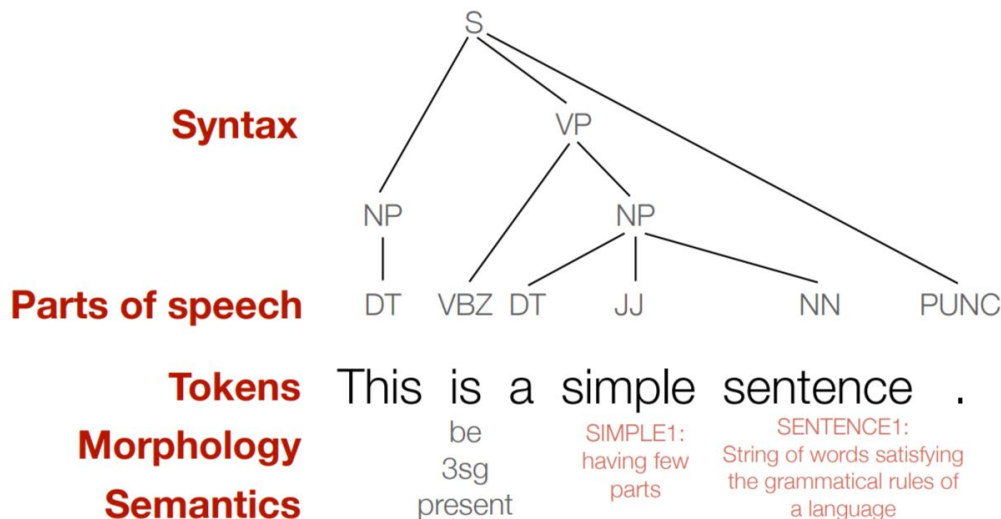
Syntax

- Part-of-speech tagging
- Syntactic parsing



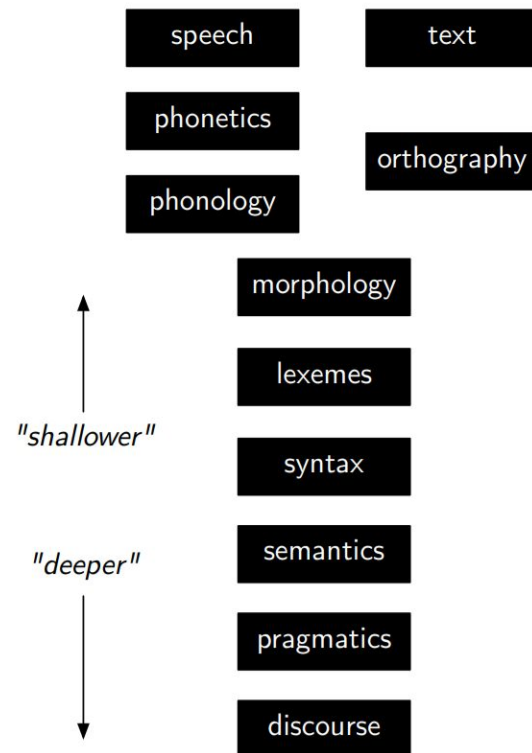
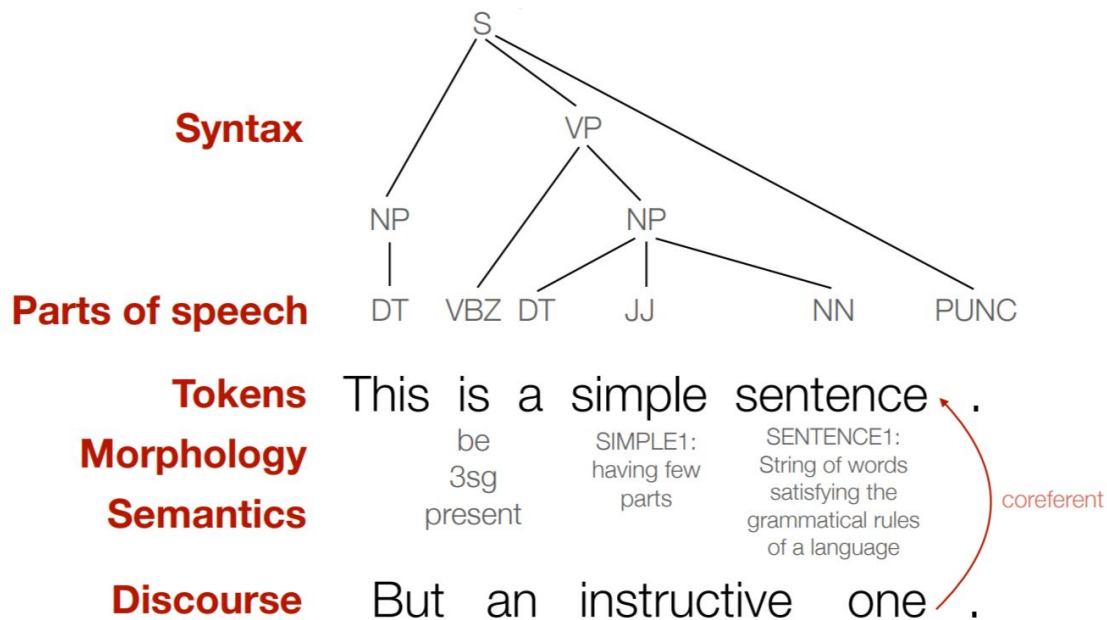
Semantics

- Named entity recognition
- Word sense disambiguation
- Semantic role labeling



Discourse

- Reference resolution
- Discourse parsing



Linguistic challenges we'll need to deal with in designing NLP systems

What are some challenges for NLP systems?

1. Ambiguity
2. Variation
3. Sparsity
4. Expressivity
5. Unmodeled variables
6. Unknown representation \mathcal{R}


Ambiguity

- Ambiguity at multiple levels:
 - Word senses: **bank** (finance or river?)
 - Part of speech: **chair** (noun or verb?)
 - Syntactic structure: **I can see a man with a telescope**
 - Multiple: **I saw her duck**



Dealing with ambiguity

- How can we model ambiguity and choose the correct analysis in context?
 - non-probabilistic methods (FSMs for morphology, CKY parsers for syntax) return *all possible analyses*.
 - probabilistic models (HMMs for part-of-speech tagging, PCFGs for syntax) and algorithms (Viterbi, probabilistic CKY) return *the best possible analysis*, i.e., the most probable one according to the model
 - Neural networks, pretrained language models now provide end-to-end solutions



- But the “best” analysis is only good if our probabilities are accurate. Where do they come from?

Corpora

- A corpus is a collection of text
 - Often annotated in some way
 - Sometimes just lots of text
- Examples
 - Penn Treebank: 1M words of parsed WSJ
 - Canadian Hansards: 10M+ words of aligned French / English sentences
 - Yelp reviews
 - The Web: billions of words of who knows what



What are some challenges for NLP systems?

1. Ambiguity
2. Variation
3. Sparsity
4. Expressivity
5. Unmodeled variables
6. Unknown representation \mathcal{R}

Scale

- ~7K languages
- Thousands of language varieties



Englishes



Africa is a continent with a very high linguistic diversity: there are an estimated 1.5-2K African languages from 6 language families. **1.33 billion people**

NLP beyond English

- ~7,000 languages
- thousands of language varieties

AT&T 6:56 PM

"Necesito una reserva para cenar para el día de San Valentín"

Veré si algún restaurante tiene una mesa para uno.

"No. Necesito una reserva para dos."

¿Por qué? ¿Está tu madre en la ciudad?

Spanish
534 million speakers

AT&T 6:56 PM

"मुझे वेल्फेडिन डे के लिए रात के खाने के आरक्षण की आवश्यकता है"

मैं देखूंगा कि क्या किसी रेस्तरां में एक के लिए एक तालिका है।

"नहीं, मुझे दो आरक्षण चाहिए।"

क्यों? क्या तुम्हारी माँ शहर में है?

Hindi
615 million speakers

AT&T 6:56 PM

"Nahitaji uhifadhi wa chakula cha jioni kwa siku ya wapendanao"

Nitaona ikiwa mikahawa yoyote inayo meza moja.

"Hapana. Ninahitaji uhifadhi wa mbili."

Kwa nini? Je! Mama yako yuko mjini?

Swahili
100 million speakers

AT&T 6:56 PM

"I need a dinner reservation for Valentine's day"

I'll see if any restaurants have a table for one.

"No. I need a reservation for two."

Why? Is your mother in town?

American English

AT&T 6:56 PM

"Ah need a tatties an' neeb's reservation fur Valentine's day ."

I'll see if onie restaurants hae a table fur a body.

"Nae. Ah need a reservation fur tois."

Wa? is yer maw in toon?

Scottish English

AT&T 6:56 PM

"Mujhe Valentine's day par reservation chahiye."

I'll see agar ek aadmi ke liye table hai.

"Nhi. Mujhe do logo ke liye table chahiye."


Kyu? Aapki mother town me hain?

Hinglish

Most of the world today is multilingual


Percentage of Bilingual Speakers in the World

European Union



Source: European Commission, "Europeans and their Languages," 2006

Percentage of US Population who spoke a language other than English at home by year




Source: U.S. Census Bureau, 2007 American Community Survey

Source: US Census Bureau

The Countries With The Most Spoken Languages


Number of living languages spoken per country in 2015



Source: Ethnologue

Semantic analysis

- Every language represents the world in a different way
 - For example, it could depend on cultural or historical conditions



- Russian has very few words for colors, Japanese has hundreds
- Multiword expressions, e.g. **happy as a clam**, **it's raining cats and dogs** or **wake up** and metaphors, e.g. **love is a journey** are very different across languages

Tokenization

这是一个简单的句子

WORDS

This is a simple sentence

זה משפט פשוט

Tokenization + disambiguation

in tea
her daughter

בתה

in tea
in the tea
that in tea
that in the tea
and that in the tea

בתה
בהתה
שבתה
שבהתה
ושבהתה

ושבתה

- most of the vowels unspecified

and her saturday
and that in tea
and that her daughter

ו+שבת+ה
ו+ש+ב+ה+תה
ו+ש+בת+ה

- most of the vowels unspecified
- particles, prepositions, the definite article, conjunctions attach to the words which follow them
- tokenization is highly ambiguous

Tokenization + morphological analysis

- Quechua

Much'anamayakapushasqakupuniñataqsunamá


Much'a -na -naya -ka -pu -sha -sqa -ku -puni -ña -taq -suna -má

"So they really always have been kissing each other then"

Much'a	to kiss
-na	expresses obligation, lost in translation
-naya	expresses desire
-ka	diminutive
-pu	reflexive (kiss *eachother*)
-sha	progressive (kiss*ing*)
-sqa	declaring something the speaker has not personally witnessed
-ku	3rd person plural (they kiss)
-puni	definitive (really*)
-ña	always
-taq	statement of contrast (...then)
-suna	expressing uncertainty (So...)
-má	expressing that the speaker is surprised


Tokenization + morphological analysis

- German



Infektionsschutzmaßnahmenverordnung

NLP Technologies/Applications



Linguistic variation

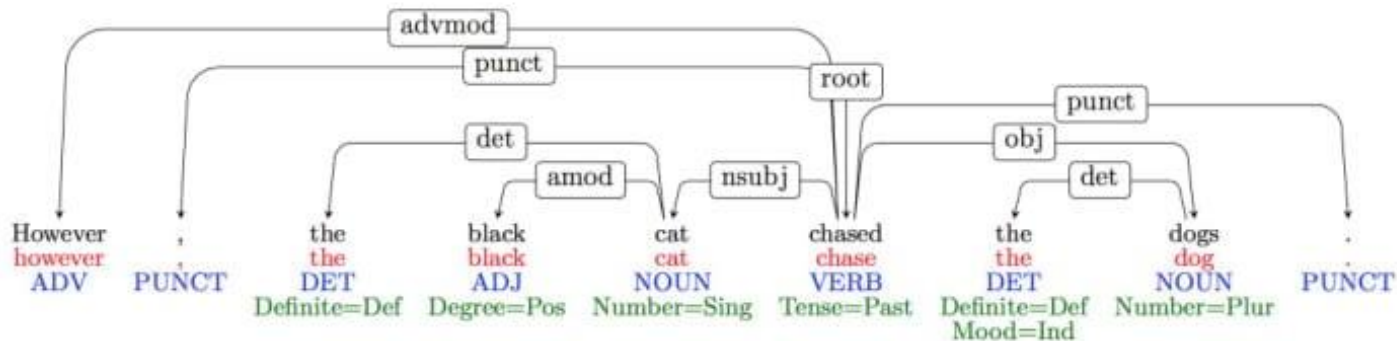
- Non-standard language, emojis, hashtags, names



chowdownwithchan #crab and #pork #xiaolongbao at @dintaifungusa... where else? 🤔👩🏻👦 Note the cute little crab indicator in the 2nd pic 🦀💕💕

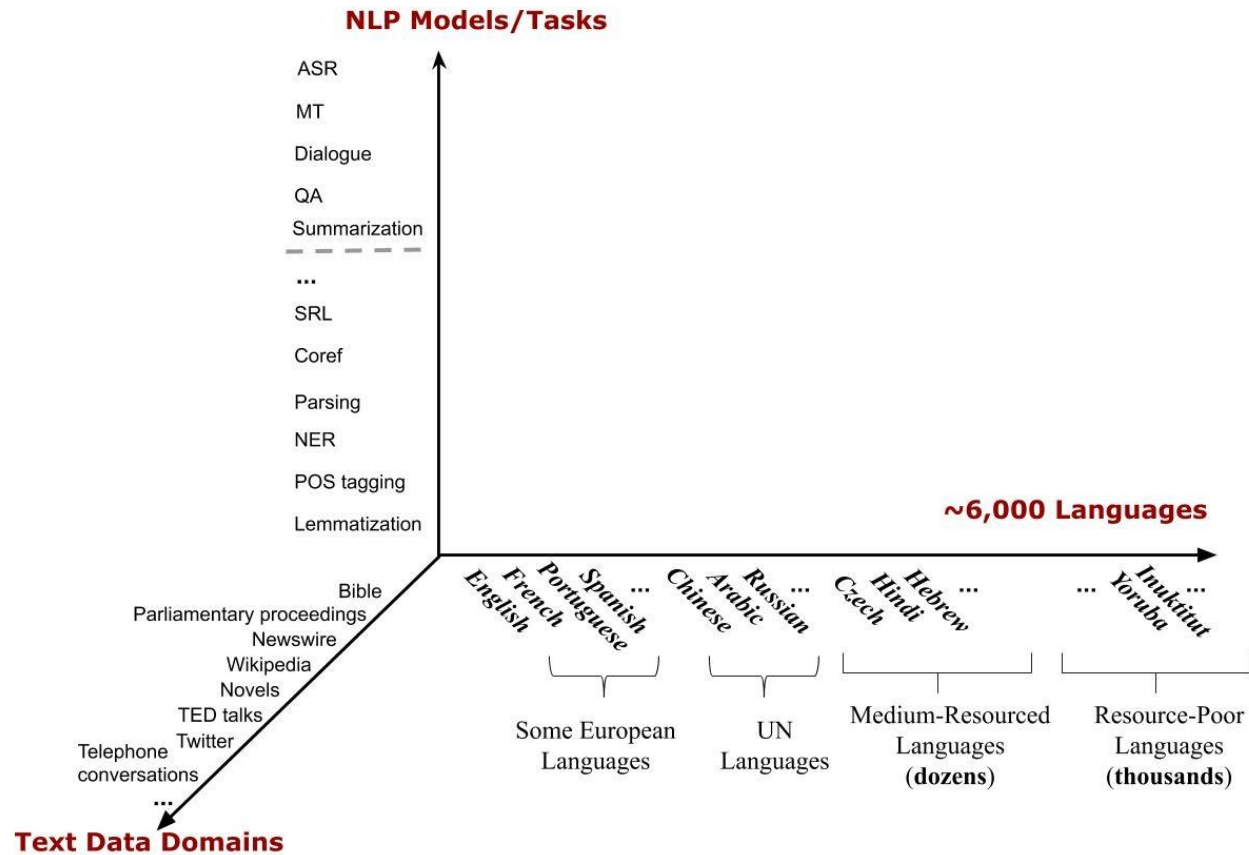
Variation

- Suppose we train a part of speech tagger or a parser on the Wall Street Journal



- What will happen if we try to use this tagger/parser for social media??

@_rkpntrnte hindi ko alam babe eh, absent ako kanina I'm sick rn hahaha 🙄🙌



What are some challenges for NLP systems?

1. Ambiguity
2. Variation
3. Sparsity
4. Expressivity
5. Unmodeled variables
6. Unknown representation \mathcal{R}

Sparsity

Sparse data due to **Zipf's Law**

- To illustrate, let's look at the frequencies of different words in a large text corpus
- Assume “word” is a string of letters separated by spaces

Word Counts

Most frequent words in the English Europarl corpus (out of 24m word tokens)

any word		nouns	
Frequency	Token	Frequency	Token
1,698,599	the	124,598	European
849,256	of	104,325	Mr
793,731	to	92,195	Commission
640,257	and	66,781	President
508,560	in	62,867	Parliament
407,638	that	57,804	Union
400,467	is	53,683	report
394,778	a	53,547	Council
263,040	I	45,842	States

Word Counts



But also, out of 93,638 distinct words (word types), 36,231 occur only once.

Examples:

- cornflakes, mathematicians, fuzziness, jumbling
- pseudo-rapporteur, lobby-ridden, perfunctorily,
- Lycketoft, UNCITRAL, H-0695
- policyfor, Commissioneris, 145.95, 27a

Plotting word frequencies





Order words by frequency. What is the frequency of nth ranked word?



Zipf's Law

Implications

- Regardless of how large our corpus is, there will be a lot of infrequent (and zero-frequency!) words
- This means we need to find clever ways to estimate probabilities for things we have rarely or never seen



What are some challenges for NLP systems?

1. Ambiguity
2. Variation
3. Sparsity
4. Expressivity
5. Unmodeled variables
6. Unknown representation \mathcal{R}

Expressivity

Not only can one form have different meanings (ambiguity) but the same meaning can be expressed with different forms:

She gave the book to Tom vs. She gave Tom the book

Some kids popped by vs. A few children visited

Is that window still open? vs. Please close the window


What are some challenges for NLP systems?

1. Ambiguity
2. Variation
3. Sparsity
4. Expressivity
5. **Unmodeled variables**
6. Unknown representation \mathcal{R}

Unmodeled variables



“Drink this milk”



World knowledge

- I dropped the glass on the floor and it broke
- I dropped the hammer on the glass and it broke

What are some challenges for NLP systems?

1. Ambiguity
2. Variation
3. Sparsity
4. Expressivity
5. Unmodeled variables
6. Unknown representation \mathcal{R}

Unknown representation

- Very difficult to decide on a representation \mathcal{R} , since we don't even know how to represent the knowledge a human has/needs:
 - What is the “meaning” of a word or sentence?
 - How to model context?
 - Other general knowledge?

Desiderata for NLP models

- Sensitivity to a wide range of phenomena and constraints in human language
- Generality across languages, modalities, genres, styles
- Computational efficiency at construction time and runtime
- Strong formal guarantees (e.g., convergence, statistical efficiency, consistency)
- High accuracy when judged against expert annotations and/or test data specific to a particular task
- Explainable to human users
- Ethical

Next class

- Text classification

Questions?